

# LLM Foundations for Agents

A CRASH COURSE ON THE CORE MECHANICS OF LARGE LANGUAGE MODELS AS THEY RELATE TO AUTONOMOUS AGENTS | LEVEL:  
BEGINNER TO INTERMEDIATE | ESTIMATED DURATION: 60 MINUTES

BY AI Course Designer

2023-10-27

---

# Course Roadmap

- [ ] Understanding Transformer Architecture: The foundation of modern LLMs.
- [ ] Mastering Prompt Engineering for Reasoning: Moving beyond simple queries to complex logic.
- [ ] Exploring Memory Systems (RAG): Giving AI a long-term memory and external knowledge.
- [ ] Learning Tool-Use Capabilities: Enabling agentic behavior through API and tool integration.

# The Architecture of LLMs

A deep dive into the evolution of neural networks and the structural breakthroughs that make modern Large Language Models possible.

---

# The Transformer Engine

The Transformer architecture shifts away from sequential processing (like RNNs) to a **parallelized approach** using the Attention mechanism. This allows the model to weigh the importance of different parts of the input data regardless of their distance in the sequence.


**Self-Attention:** Enables a token to look at other tokens in the same sequence to better understand its own context.

**Multi-Head Attention:** Parallelizes the attention process, allowing the model to attend to information from **different representation subspaces** simultaneously.

**Positional Encoding:** Since the model lacks recurrence, these vectors are added to inputs to provide information about the **relative position** of tokens.

# The Transformer Engine

MERMAID DIAGRAM



High-level architecture of the Transformer Encoder-Decoder structure

---

# Tokens, Embeddings, and Latent Space

Large Language Models (LLMs) cannot process raw text; they require numerical input. This transformation happens in three critical stages: **Tokenization**, **Embedding**, and positioning within the **Latent Space**.

**Tokens:** The basic building blocks of text. A token can be a whole word, a part of a word (subword), or a single character.

**Embeddings:** High-dimensional vectors (lists of numbers) that represent the semantic meaning of a token. Words with similar meanings have similar vectors.

**Latent Space:** The multi-dimensional mathematical space where embeddings reside. Distance and direction in this space represent conceptual relationships.

# Tokens, Embeddings, and Latent Space

| METHOD                   | APPROACH                                                        | KEY ADVANTAGE                                             | EFFICIENCY IMPACT                                            |
|--------------------------|-----------------------------------------------------------------|-----------------------------------------------------------|--------------------------------------------------------------|
| BPE (Byte Pair Encoding) | Merges the most frequent adjacent pairs of characters/sequences | Excellent balance between vocab size and sequence length  | Reduces 'Unknown' tokens while keeping sequences manageable  |
| WordPiece                | Merges pairs that maximize the likelihood of the training data  | Highly precise subword segmentation for language modeling | Very efficient for models like BERT; optimizes token utility |
| Character-level          | Treats every single character as an individual token            | Zero 'Out-of-Vocabulary' (OOV) issues                     | Creates very long sequences, increasing computational cost   |

# Training Paradigms

Creating a state-of-the-art Large Language Model (LLM) is not a single step, but a **multi-stage evolution**. It transforms a raw statistical predictor into a helpful, safe, and conversational assistant through a process called the **training pipeline**.

**Pre-training:** The foundation phase where the model learns general language patterns from massive, diverse corpora.

**Supervised Fine-Tuning (SFT):** The instruction phase where the model learns to follow specific prompts using curated (Input, Output) pairs.

**RLHF:** The alignment phase using **Reinforcement Learning from Human Feedback** to refine outputs based on human preferences.

# Prompt Engineering for Reasoning

Mastering the techniques to guide LLMs through complex, multi-step logic and reasoning processes essential for building autonomous agentic workflows.

---

# Prompting Strategies

Prompting is the art of communicating with an AI to get the **best possible result**. Depending on the complexity of the task, you can use different strategies to guide the model's reasoning and output format.

**Zero-shot:** Asking the model to perform a task without any prior examples.

**Few-shot:** Providing a few examples to help the model understand the **pattern or format** required.

**Chain-of-Thought (CoT):** Encouraging the model to show its **step-by-step reasoning** before giving the final answer.

# Prompting Strategies

| STRATEGY         | WHAT IS IT?                 | BEST USE CASE                        | REASONING LEVEL |
|------------------|-----------------------------|--------------------------------------|-----------------|
| Zero-shot        | No examples provided        | Simple tasks or general knowledge    | Basic           |
| Few-shot         | A few input-output examples | Specific formats or pattern matching | Intermediate    |
| Chain-of-Thought | Step-by-step logic guidance | Complex math or multi-step logic     | Advanced        |

---

# Advanced Reasoning Frameworks

To move beyond simple linear prompting, advanced frameworks allow LLMs to simulate human-like problem-solving. **Tree-of-Thoughts (ToT)** enables the model to explore multiple reasoning paths simultaneously, while **ReAct** integrates reasoning with external tool usage.


**Tree-of-Thoughts (ToT):** Solves complex problems by branching out multiple 'thought' paths and evaluating which one is most promising.

**ReAct (Reason + Act):** An interleaved pattern where the model generates a reasoning trace and a task-specific action in a loop.

**Goal:** To reduce *hallucinations* and improve the reliability of multi-step logic.

# Advanced Reasoning Frameworks

MERMAID DIAGRAM



*The ReAct (Reason + Act) Iterative Loop*

---

# The Power of System Prompts

The **System Prompt** (or System Message) is the foundational set of instructions that defines how an AI agent behaves. It acts as the 'operational manual' that the model references throughout the entire conversation to ensure consistency and reliability.

**Persona Definition:** Establishes *who* the AI is (e.g., 'You are a senior cybersecurity analyst').

**Operational Constraints:** Defines the *boundaries* of behavior (e.g., 'Do not mention competitor products').

**Response Guidelines:** Dictates the format and tone (e.g., 'Keep responses concise and use technical language').

"

*Precision in the system prompt is the difference between an AI that simply answers and an AI that solves. Steering is the most critical lever in agent engineering.*

"

– INDUSTRY EXPERT IN LLM ORCHESTRATION

# LLMs as Agentic Controllers

Understanding the transition of Large Language Models from passive chatbots to the central reasoning engines of autonomous agents.

---

# Tool Use and Function Calling

Large Language Models (LLMs) are incredible at reasoning, but they have inherent limitations: they lack **real-time data**, cannot perform **complex calculations** reliably, and cannot interact with **external APIs** on their own. **Function Calling** is the mechanism that allows an LLM to act as a controller, identifying when a specific tool is needed and requesting its execution.


**Bridging the Gap:** Connects static model weights to **dynamic, real-time data**.

**Structured Output:** The LLM outputs **JSON arguments** instead of conversational text.

**Deterministic Actions:** Enables LLMs to trigger actual code execution in a controlled environment.

# Tool Use and Function Calling

MERMAID DIAGRAM



*The Function Calling Lifecycle: From Intent to Execution*

---

# Planning and Decomposition

In the world of AI agents, **Planning** is the process of determining a sequence of actions to achieve a specific goal. Because complex goals (e.g., "Research and write a market report") are too broad for a single LLM prompt, agents use **Decomposition** to break these macro-goals into smaller, manageable **sub-tasks**.

**Goal Decomposition:** The act of splitting a high-level objective into discrete, executable steps.

**Task Dependency:** Identifying which sub-tasks must be completed *before* others can begin.

**Iterative Refinement:** The ability of an agent to re-evaluate and adjust the plan based on the outcome of previous steps.

# Planning and Decomposition

| FEATURE       | SEQUENTIAL PLANNING                               | HIERARCHICAL PLANNING                          |
|---------------|---------------------------------------------------|------------------------------------------------|
| Structure     | Linear (Step A → Step B → Step C)                 | Nested (Goal → Sub-goal → Task)                |
| Complexity    | Low; easy to implement and track                  | High; requires recursive logic                 |
| Flexibility   | Rigid; failure in one step often stalls the chain | Flexible; can pivot within a specific sub-goal |
| Best Use Case | Simple, predictable workflows                     | Complex, multi-stage projects                  |

---

# Agentic Loops and Self-Correction

An **Agentic Loop** is a design pattern where an LLM moves beyond 'one-shot' generation. Instead of providing a single final answer, the agent enters a cycle of **self-reflection**, acting as both the *creator* and the *critic* to iteratively refine its output until it reaches a correct or optimized solution.

**Initial Generation:** The agent produces a first-draft solution based on the prompt.

**Self-Critique:** The agent reviews its own output for errors, hallucinations, or missing constraints.

**Iterative Refinement:** The agent uses the critique to rewrite and correct the previous output.

**Termination:** The loop exits once the agent validates the answer or reaches a maximum iteration limit.

# Memory and Context Management

Exploring how AI agents maintain state and retrieve relevant information over long-term interactions to ensure coherence and personalization.

# Short-term vs Long-term Memory

In the world of Large Language Models, 'memory' is handled differently than in humans. It is split between the **Context Window** (immediate processing) and **External Databases** (persistent storage).

**Short-term Memory (Context Window):** The limited amount of text the model can 'see' and consider at any single moment during a conversation.

**Long-term Memory (External Databases):** Using techniques like **RAG (Retrieval-Augmented Generation)** to fetch relevant data from a vector database and feed it into the context window.

**The Core Limitation:** As the context window grows, computational cost increases and the model may suffer from the '*lost-in-the-middle*' phenomenon.

# Short-term vs Long-term Memory

| MODEL             | PROVIDER  | CONTEXT WINDOW (TOKENS) |
|-------------------|-----------|-------------------------|
| GPT-4o            | OpenAI    | 128,000                 |
| Claude 3.5 Sonnet | Anthropic | 200,000                 |
| Gemini 1.5 Pro    | Google    | 2,000,000               |

---

# Retrieval Augmented Generation (RAG)

Retrieval Augmented Generation, or **RAG**, is a powerful architecture that optimizes the output of a Large Language Model (LLM) by referencing a specific, authoritative knowledge base outside of its original training data. Essentially, it turns the LLM from a closed-book exam taker into an **open-book** researcher.

**Reduces Hallucinations:** Grounds the model in factual, retrieved data.

**Up-to-Date Information:** Provides access to real-time data without requiring expensive **model retraining**.

**Domain Specificity:** Allows the AI to answer questions about *private or proprietary* documents securely.

# Retrieval Augmented Generation (RAG)

MERMAID DIAGRAM



*The end-to-end RAG Pipeline flow*

# Vector Databases and Semantics

Unlike traditional databases that rely on exact keyword matches, **Vector Databases** store information as **embeddings**—numerical representations of meaning in a high-dimensional space. This shift allows AI systems to perform **semantic search**, identifying content based on its *conceptual meaning* rather than just the words used.

**Embeddings:** Transforming text, images, or audio into vectors (long lists of numbers) using a model.

**Semantic Space:** Concepts with similar meanings are positioned **mathematically close** to each other in a vector space.

**Knowledge Retrieval:** This architecture allows LLMs to access external, real-time data, a process known as **Retrieval-Augmented Generation (RAG)**.

# Reliability and Guardrails

Exploring the challenges of deploying LLM agents in production, focusing on safety, predictability, and stability.

---

# Managing Hallucinations

An LLM hallucination occurs when the model generates text that is **factually incorrect** or nonsensical, despite appearing confident and fluent. This happens because LLMs are **probabilistic engines**, not databases; they predict the next most likely token based on patterns rather than retrieving verified facts.

**Probabilistic Nature:** Models prioritize fluency and likelihood over factual accuracy.

**Data Gaps:** Hallucinations often arise when the training data is *incomplete or outdated*.

**Over-optimization:** The drive to be helpful can lead the model to 'fill in the blanks' creatively.

# Managing Hallucinations

- ☒ Grounding via RAG (Retrieval-Augmented Generation)
- ☒ Temperature Control (Reducing randomness)
- ☐ Multi-step Verification & Self-Correction
- ☐ Few-shot Prompting for Format Constraints

# Output Structuring and Validation

LLMs are natively conversational, but for programmatic integration, we need **deterministic structured data**. Relying on raw text prompts often leads to 'hallucinated' formats or conversational filler that breaks downstream parsers. To solve this, we move from simple prompting to **Schema Enforcement**.

**The Challenge:** LLMs are probabilistic, meaning they may omit brackets or add preamble text.

**JSON Mode & Function Calling:** Native API features that force the model to output valid JSON.

**Pydantic Validation:** Using Python classes to define strict types, required fields, and value constraints.

**Guardrail Frameworks:** Using tools like *Guardrails AI* or *Outlines* to validate and auto-correct outputs in real-time.

---

## Safety and Ethics in Agency

Giving AI agents autonomous access to tools increases efficiency but introduces **significant safety risks**. Without constraints, an agent might misinterpret a command and perform irreversible actions, such as deleting critical files or sending unauthorized communications.

**Unintended Actions:** The risk of 'hallucinating' a tool call or executing a destructive command.

**Security Vulnerabilities:** Potential for agents to be manipulated via prompt injection to bypass safety filters.

**Human-in-the-Loop (HITL):** A critical checkpoint where the agent must seek **explicit human approval** before executing high-risk operations.

"

*The goal of AI alignment is to ensure that the agent's goals are perfectly aligned with human values and intentions, preventing the gap between what we ask for and what we actually want.*

"

– AI ALIGNMENT RESEARCH PRINCIPLES

---

# Course Summary and Next Steps

Course Wrap-up: Mastering LLM Foundations

Key Pillars Summary:

Architecture: Understanding Transformers and how LLMs process language.

Prompting: Mastering the art of steering model output for precision.

Tools: Enabling LLMs to interact with the real world via APIs and functions.

Memory: Implementing RAG and context management for persistence.

Safety: Applying guardrails and alignment to ensure ethical, reliable AI.

 Your Next Step: Build Your First Agent! Combine these five pillars to create an autonomous agent that can plan, remember, and execute tasks.

Thank you for completing the course!